

OneInbox Agents Guide v1

OneInbox Agents - Concept Guide

What is an Agent?

An agent is the voice AI your callers talk to. It has a personality, a voice, and a set of actions it can take during a live phone or web conversation.

Every agent has four building blocks:

1. Voice - what the agent sounds like (Cartesia, Deepgram, ElevenLabs, OpenAI)
2. LLM - the AI model that powers the agent's reasoning (GPT-4o, Claude, Gemini)
3. Tools - actions the agent can take mid-call (send SMS, transfer call, extract info)
4. Knowledge Base - documents the agent can reference to answer questions

Creating an Agent

To create an agent, make a POST request to the OneInbox API:

POST <https://api-tokyo.oneinbox.ai/v1/agents>

Headers: Authorization: Bearer <api_key>

Content-Type: application/json

Request body:

```
{
  "name": "Support Agent",
  "system_prompt": "You are a helpful customer support agent for Acme Corp.",
  "voice_id": "<voice_id>",
  "llm_id": "<llm_id>"
}
```

The system prompt is how you control the agent's personality and behavior. Write it as instructions the agent follows during every call. Reference tools by their name field.

Agent Configuration Options:

- name: Display name for the agent
- system_prompt: Instructions the agent follows during calls
- voice_id: UUID of the voice to use
- llm_id: UUID of the LLM model configuration
- knowledge_base_ids: List of KB UUIDs for document retrieval
- first_message: What the agent says when the call connects
- language: Primary language (default: en)

How a Call Works

When a call connects, the following pipeline runs in real time:

Caller speaks -> speech is transcribed to text -> LLM decides what to say -> response is spoken back to the caller.

OneInbox handles real-time audio streaming, speech recognition, telephony routing, and scaling automatically. The entire loop runs with sub-second latency.

Authentication

OneInbox uses two types of API keys:

- API Key (starts with oi_sk_): Full access, used server-side
- Publishable Key (starts with oi_pk_): Browser-only, can only start calls

Use your API key for all server-side operations. The publishable key is for the Web SDK only.

What You Can Build

- Outbound calling bots: Agent dials leads, qualifies them, books demos
- Inbound support lines: Assign a phone number, callers reach your agent 24/7
- Website voice widgets: Web SDK lets visitors talk to your agent in the browser